

Saqib Rashid Bhat

6005370241 | saqib70241@gmail.com | LinkedIn | Portfolio | github.com/Saqibbhat321

Professional Summary

Aspiring Machine Learning Engineer with hands-on experience building end-to-end AI and ML applications using **Python, FastAPI, Scikit-Learn, FAISS, MLflow, Docker, and LLMs**. Developed projects across **RAG, semantic search, recommendation, and predictive modeling**, with practical exposure to **NLP, Retrieval-Augmented Generation (RAG), Information Retrieval, model evaluation, REST API development, and MLOps**. Interested in building reliable ML systems that move from experimentation to deployable applications.

Technical Skills

Languages: Python, SQL, Java

ML & Data Science: Scikit-Learn, XGBoost, Random Forest, Logistic Regression, Classification, Regression, Feature Engineering, Cross-Validation, Hyperparameter Tuning, Model Evaluation, ROC-AUC, Statistical Learning

NLP / GenAI / RAG: Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), Natural Language Processing (NLP), Prompt Engineering, Semantic Search, Sentence Transformers, Dense Embeddings, Similarity Search, Information Retrieval

Data & Visualization: Pandas, NumPy, Matplotlib, Seaborn, EDA, Data Preprocessing, Data Wrangling, Feature Extraction

Backend / Databases / MLOps: FastAPI, Pydantic, PostgreSQL, MySQL, FAISS, SQLAlchemy, Alembic, MLflow, Docker, Docker Compose, Streamlit, Git, GitHub, pytest, Jupyter

Projects

Intelligent Document Q&A Assistant (RAG)

GitHub

Python, FastAPI, RAG, FAISS, Sentence Transformers, Groq (Llama-3.3-70B), PostgreSQL, MLflow, Streamlit, Docker

- Built a Retrieval-Augmented Generation (RAG) application that ingests **PDF, DOCX, and TXT** documents and answers user questions with **source-grounded responses** by combining semantic retrieval, LLM generation, and document-level citation support.
- Engineered the retrieval pipeline using **Sentence Transformers (all-MiniLM-L6-v2)** and a **FAISS** vector index to support semantic cosine-similarity search across mixed-format document collections.
- Integrated the **Groq-hosted Llama-3.3-70B** model via an OpenAI-compatible API to generate context-aware answers from retrieved passages with source attribution for improved answer traceability.
- Implemented per-session conversation memory in **PostgreSQL** and logged retrieval latency, response latency, and experiment metadata with **MLflow** to support iterative tuning of the RAG pipeline.
- Delivered a **FastAPI** backend with **4 REST endpoints** and a **Streamlit** chat UI, fully containerized with **Docker Compose** across **4 services**.

AI Resume & Job-Match Assistant

GitHub

Python, FastAPI, FAISS, Sentence Transformers, LLM, SQLAlchemy, Docker

- Built an AI job-matching assistant using **Sentence-Transformer embeddings, cosine similarity**, and a **FAISS** vector index to rank jobs beyond exact keyword matching.
- Parsed resumes (**PDF/DOCX**), extracted skills, computed **ATS-style skill-gap scores**, and generated LLM-powered resume suggestions with a rule-based fallback.
- Developed a service-oriented **FastAPI** backend with **SQLAlchemy** ORM and containerized the application with **Docker**.

Customer Churn Prediction System

GitHub

Python, Scikit-Learn, FastAPI, Docker, MLflow, Pandas, NumPy, Joblib, Pydantic

- Built an end-to-end **Customer Churn Prediction System** using a **Random Forest Classifier** on telecom customer data, achieving **0.84 ROC-AUC** and **80% accuracy** through data preprocessing, feature engineering, model training, and evaluation.
- Engineered a reusable machine learning pipeline using **ColumnTransformer, OneHotEncoder, and StandardScaler** to ensure consistent preprocessing across both **training and inference**.
- Deployed the model as a **FastAPI REST API** with **Pydantic** validation, **Docker** containerization, prediction logging, and **MLflow** experiment tracking to support reproducible model serving.

Experience

Herways Coffee

June 2025 – July 2026

Operations Manager – Data, Reporting & Process Optimization

Bangalore, India

- Managed inventory, billing, vendor coordination, and operational reporting using **Google Sheets** and **Swipe**, maintaining structured sales and stock data to support day-to-day planning and process execution.
- Prepared recurring sales, inventory, and vendor transaction reports and analyzed stock movement trends to support replenishment planning, improve record accuracy, and strengthen operational visibility.

Education

Cambridge Institute of Technology

2020 – 2024

Bachelor of Engineering in Computer Science

Bangalore, Karnataka

Minor in Artificial Intelligence & Data Science (IIT Mandi × Masai School)

2025 – 2026

Machine Learning, Natural Language Processing, Deep Learning, Probability & Statistics, Linear Algebra, Data Analytics Online