

Saqib Rashid Bhat

Machine Learning Engineer | AI / LLM & RAG Systems

+91-6005370241 | saqib.bhat.x@gmail.com | linkedin.com/in/saqib70241 | portfolio | github.com/Saqibbhat321

PROFESSIONAL SUMMARY

Machine Learning / AI Engineer who has shipped **3 end-to-end AI applications** a full RAG pipeline, an LLM-powered job-matching engine, and a production ML classifier (**0.84 ROC-AUC**) each built, containerized, and served through REST APIs. Hands-on with **RAG, dense embeddings, FAISS vector search, FastAPI, and MLflow**, with a track record of taking models from notebook experiments to **Docker-deployed, API-served systems**. Looking to bring RAG and backend engineering skills to an enterprise-scale AI platform team.

TECHNICAL SKILLS

Languages: Python, SQL, Java

LLM / GenAI / RAG: Retrieval-Augmented Generation (RAG), Large Language Models (LLMs), Prompt Engineering, Sentence Transformers, Dense Embeddings, Vector Search, Semantic Search, Information Retrieval, Model Orchestration (OpenAI-compatible APIs)

Backend & APIs: FastAPI, Pydantic, REST API Design, SQLAlchemy, Alembic, Microservice-style Architecture

Databases & Vector Search: PostgreSQL, MySQL, FAISS (Vector Database), SQL

ML & Data Science: Scikit-Learn, XGBoost, Random Forest, Logistic Regression, Feature Engineering, Cross-Validation, Hyperparameter Tuning, ROC-AUC Evaluation, Statistical Learning

MLOps & DevOps: Docker, Docker Compose, MLflow (Experiment Tracking), Git, GitHub

Data Engineering: Pandas, NumPy, EDA, Data Preprocessing, Data Wrangling, Feature Extraction

PROJECTS

Intelligent Document Q&A Assistant (RAG)

GitHub

Python, FastAPI, RAG, FAISS, Sentence Transformers, Groq (Llama-3.3-70B), PostgreSQL, MLflow, Streamlit, Docker

- Built a full-stack **RAG pipeline** ingesting PDF, DOCX, and TXT files to deliver **source-grounded, citation-backed answers** across mixed-format document collections.
- Engineered the retrieval layer with **Sentence Transformers (all-MiniLM-L6-v2)** and a **FAISS** vector index for sub-second cosine-similarity search over dense embeddings.
- Integrated **Llama-3.3-70B (via Groq)** through an OpenAI-compatible interface to generate context-aware, traceable answers from retrieved passages.
- Instrumented the pipeline with **MLflow**, logging retrieval and response latency to drive iterative performance tuning.
- Shipped a **FastAPI backend (4 REST endpoints)** and **Streamlit** chat UI, containerized across **4 Docker Compose services** with persistent PostgreSQL-backed session memory.

AI Resume & Job-Match Assistant

GitHub

Python, FastAPI, FAISS, Sentence Transformers, LLM, SQLAlchemy, Docker

- Built an **embedding-based job-matching engine** using Sentence-Transformer vectors and **FAISS** cosine-similarity search, ranking roles beyond simple keyword overlap.
- Automated resume parsing (PDF/DOCX) and skill extraction to generate **ATS-style skill-gap scores** and LLM-powered improvement suggestions, with a rule-based fallback for reliability.
- Delivered a service-oriented **FastAPI + SQLAlchemy** backend, fully **Dockerized** for reproducible deployment.

Customer Churn Prediction System

GitHub

Python, Scikit-Learn, FastAPI, Docker, MLflow, Pandas, NumPy, Joblib, Pydantic

- Trained a **Random Forest classifier** on telecom customer data, achieving **0.84 ROC-AUC and 80% accuracy** through systematic feature engineering and model evaluation.
- Built a reusable **ColumnTransformer-based preprocessing pipeline** (OneHotEncoder, StandardScaler) to guarantee consistent transforms across training and inference.
- Productionized the model as a **FastAPI REST service** with Pydantic validation, prediction logging, Docker containerization, and **MLflow**-tracked experiments for reproducibility.

EXPERIENCE

Herway Brew Labs

2025 – Present

Operations & Data Reporting

Bangalore, India

- Managed inventory, billing, vendor coordination, and operational reporting using Google Sheets and Swipe, maintaining structured sales and stock data to support day-to-day planning and process execution.
- Authored recurring sales, inventory, and vendor-transaction reports and analyzed stock-movement trends, improving replenishment planning and data accuracy across the business.

EDUCATION

Cambridge Institute of Technology

2020 – 2024

Bachelor of Engineering in Computer Science

Bangalore, Karnataka

Minor in Artificial Intelligence & Data Science (IIT Mandi × Masai School)

2025 – 2026

Machine Learning, Natural Language Processing, Deep Learning, Probability & Statistics, Data Analytics

Online